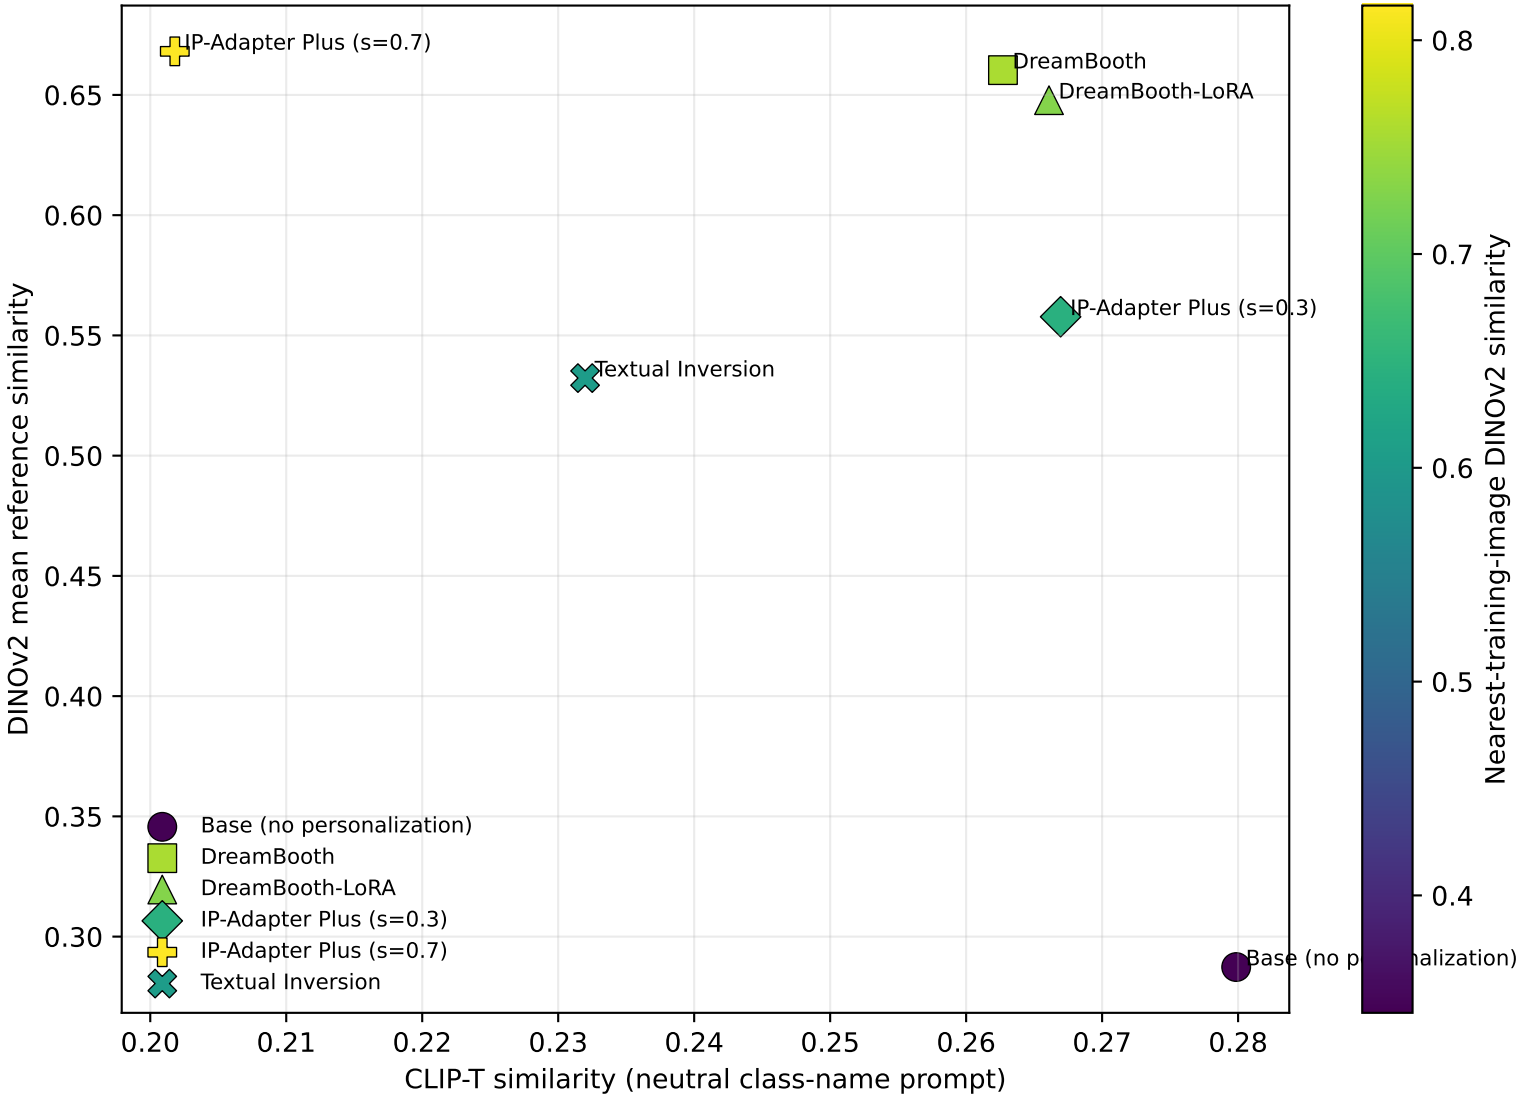
Personalization fidelity–prompt alignment trade-off
IP-Adapter Plus (s=0.7)
DreamBooth
DreamBooth-LoRA
IP-Adapter Plus (s=0.3)
Textual Inversion
Base (no personalization)
DreamBooth
DreamBooth-LoRA
IP-Adapter Plus (s=0.3)
IP-Adapter Plus (s=0.7)
Textual Inversion
DINOv2 mean reference similarity
0.65
0.60
0.55
0.50
0.45
0.40
0.35
0.30
0.20
0.21
0.22
0.23
0.24
0.25
0.26
0.27
0.28
CLIP-T similarity (neutral class-name prompt)
Nearest-training-image DINOv2 similarity
0.8
0.7
0.6
0.5
0.4